

Linear Classification Models

Perceptron

Classification problem

Let's look at the problem of spam filtering

Now anyone can learn how to earn \$200 - \$943 per day or More ! If you can type (hunt and peck is ok to start) and fill in forms, you can score big! So don't delay waiting around for the next opportunity...it is knocking now! Start here: <http://redbluecruise.com/t/c/381/polohoo/yz37957.html>

Do you Have Poetry that you think should be worth \$10,000.00 USD, we do!.. Enter our International Open contest and see if you have what it takes. To see details or to enter your own poem, Click link below. <http://e-suscriber.com/imys?e=0sAoo4q9s4zYYUoYQ&m=795314&l=0>

View my photos!
I invite you to view the following photo album(s): zak-month27

Hey have you seen my new pics yet???? Me and my girlfreind would love it if you would come chat with us for a bit.. Well join us if you interested. Join live web cam chat here: <http://e-commcentral.com/imys?e=0sAoo4q9s4zYYUoYQ&m=825314&l=0>

Spam prediction

- Training data
 - Past emails and whether they are considered spam or not (you can also choose to use non-spam or spam emails only, but that will require different choices later on)
- Target variable: spam or not
- But how do we represent the emails as features?
 - Bag of words
 - Deep learning based approach


Bag of words representation

- We need to represent each email with a fixed set of features
- Assume a fixed dictionary
- Create a feature for each dictionary word
 - Binary: whether the particular dictionary word is present in the email
 - Integer: the count of the appearances
 - Continuous: normalized count of appearances, e.g., TF-IDF (term frequency – inverse document frequency)

$$tf-idf(w_i) = \frac{\text{\# of times } w_i \text{ appears in the current email}}{\text{\# of emails in which } w_i \text{ appears}}$$

Linear Classifier

- We will begin with the simplest choice: linear classifiers




Why linear model?

- Simplest model – fewer parameters to learn (requires less training data to learn reliably)
- Intuitively appealing -- draw a straight line (for 2-d inputs) or a linear hyper-plane (for higher dimensional inputs) to separate positive from negative
- Can be used to learn nonlinear models as well. How?
 - Introducing nonlinear features (e.g., $x_1^2, x_2^2, x_1x_2 \dots$)
 - Use kernel tricks (we will talk about this later this term)

Learning Goal

- Given a set training examples, each is described by m features x_1, \dots, x_m , and belong to either the positive or negative class, $y \in \{+1, -1\}$
- Let $\mathbf{x} = [1, x_1, \dots, x_m]^T$,
- $\mathbf{w} = [w_0, w_1, \dots, w_m]^T$ defines a decision boundary $\mathbf{w}^T \mathbf{x} = 0$ that separates the input space into two parts

- The vector \mathbf{w} is the normal vector of the decision boundary, i.e., its perpendicular to the boundary
- \mathbf{w} points to the positive side
- Goal: find a \mathbf{w} s.t. the decision boundary separates positive examples from negative examples



How to learn: the perceptron algorithm

How can we achieve this? Perceptron is one approach.

1. It starts with some (random) vector \mathbf{w} and incrementally updates \mathbf{w} whenever it makes a mistake.
2. Let \mathbf{w}_t be the current weight vector, and it makes a mistake on example (\mathbf{x}, y)
 - y and $\mathbf{w}_t^T \mathbf{x}$ have different signs
 - $y\mathbf{w}_t^T \mathbf{x} < 0$
3. The update intends to correct for the mistake

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + y\mathbf{x}$$


Effect of Perceptron Updating Rule

- Mathematically speaking


$$y\mathbf{w}_{t+1}^T \mathbf{x} = y(\mathbf{w}_t + y\mathbf{x})^T \mathbf{x} = y\mathbf{w}_t^T \mathbf{x} + \underbrace{y^2 \mathbf{x}^T \mathbf{x}}_{> 0} > y\mathbf{w}_t^T \mathbf{x}$$
$$= y|\mathbf{x}|^2 > 0$$

The update makes $y\mathbf{w}_t^T \mathbf{x}$ more positive, potentially correcting the mistake

- Geometrically



Step t



Step t+1

The Perceptron Algorithm

Let $\mathbf{w} \leftarrow (0,0,0,\dots,0)$ *//Start with 0 weights*

Repeat *//go through training examples one by one*

 Accept training example $i : (\mathbf{x}_i, y_i)$

$u_i \leftarrow \mathbf{w}^T \mathbf{x}_i$ *//Apply the current weight*

 if $y_i u_i \leq 0$ *// If it is misclassified*

$\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$ *// update w*

Important notes:

- Correcting for a mistake could move the decision boundary so much that previously correct examples are now misclassified.
- As such it must go over the training examples multiple times
- Each time it goes through the whole training set, it is called an epoch.
- It will terminate if no update is made to \mathbf{w} during one epoch – which means it has converged.

Online vs Batch

- We call the above perceptron algorithm an ***online algorithm***
- Online algorithms perform learning each time it receives an training example
- In contrast, ***batch learning*** algorithms collect a batch of training examples and learn from them all at once.

Batch Perceptron Algorithm

```
Given : training examples  $(\mathbf{x}_i, y_i)$ ,  $i = 1, \dots, N$   
Let  $\mathbf{w} \leftarrow (0, 0, 0, \dots, 0)$   
do  
     $\mathbf{delta} \leftarrow (0, 0, 0, \dots, 0)$   
    for  $i = 1$  to  $N$  do  
         $u_i \leftarrow \mathbf{w}^T \mathbf{x}_i$   
        if  $y_i \cdot u_i \leq 0$   
             $\mathbf{delta} \leftarrow \mathbf{delta} - y_i \mathbf{x}_i$   
     $\mathbf{delta} \leftarrow \mathbf{delta} / N$   
     $\mathbf{w} \leftarrow \mathbf{w} - \eta \mathbf{delta}$   
until  $|\mathbf{delta}| < \varepsilon$ 
```

- Perceptron does gradient descent to minimize loss function $E(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N (-y_i \mathbf{w}^T \mathbf{x}_i)_+$
- \mathbf{delta} stores the gradient
- η is the learning rate of the gradient descent steps
- Too large η causes oscillation, too small leads to slow convergence
- Common to use large η first, then gradually reduce it

Good news

- **Convergence Property:**

For linearly separable data (i.e., there exists an linear decision boundary that perfectly separates positive and negative training examples), the perceptron algorithm converges in a **finite number of steps**.

- *Why?* If you are mathematically curious, read the following slide, you will find the answer.
- *And how many steps?* If you are practically curious, read the following slide, answer is in there too.
- **The further good news** is that you are not required to master this material, they are just provided for the curious ones

To show convergence, we just need to show that each update moves the weight vector closer to a solution vector by a lower bounded amount
Let w^* be a solution vector, and w_t be our w at t th step,

Proof

$$\text{cosine}(w^*, w_t) = \frac{w^* \cdot w_t}{\|w^*\| \cdot \|w_t\|}$$

$$w^* \cdot w_t = w^* \cdot (w_{t-1} + y^t x^t) = w^* \cdot w_{t-1} + w^* y^t x^t$$

Assume that w^* classify all examples with a margin γ , i.e., $w^* y x > \gamma$ for all examples

$$w^* \cdot w_t = w^* \cdot w_{t-1} + w^* y^t x^t > w^* \cdot w_{t-1} + \gamma > w^* \cdot w_{t-2} + 2\gamma > \dots > w^* \cdot w_0 + t\gamma = t\gamma$$


$$\|w_t\|^2 = \|w_{t-1} + y^t x^t\|^2 = \|w_{t-1}\|^2 + y^{t2} \|x^t\|^2 + 2w_{t-1} y^t x^t < \|w_{t-1}\|^2 + \|x^t\|^2$$

Assume that $\|x\|$ are bounded by D


$$\|w_t\|^2 < \|w_{t-1}\|^2 + \|x^t\|^2 < \|w_{t-1}\|^2 + D^2 < \|w_{t-2}\|^2 + 2D^2 < \dots < tD^2$$

$$\text{cosine}(w^*, w_t) = \frac{w^* \cdot w_t}{\|w^*\| \cdot \|w_t\|} > \frac{t\gamma}{\|w^*\| \cdot \|w_t\|} > \frac{t\gamma}{\|w^*\| \cdot \sqrt{tD^2}}$$

$$\frac{t\gamma}{\|w^*\| \cdot \sqrt{tD^2}} < 1 \Rightarrow \sqrt{t} < \frac{D\|w^*\|}{\gamma} \Rightarrow t < D^2 \left(\frac{\gamma^2}{\|w^*\|^2} \right)$$



Bad news



What about non-linearly separable cases!

In such cases the algorithm will never stop! How to fix?

One possible solution: look for decision boundary that make as few mistakes as possible – NP-hard (refresh your CS325 memory!)

Fixing the Perceptron

Idea one: only go through the data once, or a fixed number of times

```
Let  $\mathbf{w} \leftarrow (0,0,0,\dots,0)$ 
for  $i = 1,\dots,T$ 
    Take training example  $i : (\mathbf{x}_i, y_i)$ 
     $u_i \leftarrow \mathbf{w}^T \mathbf{x}_i$ 
    if  $y_i u_i \leq 0$ 
         $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$ 
```

At least this stops!

Problem: the final \mathbf{w} might not be good
e.g., right before terminating, the algorithm might perform an update on a total outlier...

Voted-Perceptron

Idea two: keep around intermediate hypotheses, and have them “vote” [Freund and Schapire, 1998]

```
Let  $w_0 = (0,0,0,...,0)$ 
 $c_0 = 0, n = 0$ 
repeat for a fixed number of steps
  Take example  $i : (\mathbf{x}_i, y_i)$ 
   $u_i \leftarrow \mathbf{w}_n^T \mathbf{x}_i$ 
  if  $y_i u_i \leq 0$ 
     $\mathbf{w}_{n+1} \leftarrow \mathbf{w}_n + y_i \mathbf{x}_i$ 
     $c_{n+1} = 0$ 
     $n = n + 1$ 
  else
     $c_n = c_n + 1$ 
```

Store a collection of linear separators w_0, w_1, \dots , along with their survival time c_0, c_1, \dots

The c 's can be good measures of reliability of the w 's.

For classification, take a weighted vote among all N separators:

$$\text{sgn} \left\{ \sum_{n=0}^N c_n \text{sgn}(\mathbf{x}^T \mathbf{w}_n) \right\}$$

Average perceptron:

$$\mathbf{w}_A = \sum_{n=0}^N c_n \mathbf{w}_n$$

Summary

- Perceptron incrementally learns a linear decision boundary to separate positive from negative
- It begins with a random weight vector or a zero weight vector, and incrementally update the weight vector whenever it makes a mistake
- Each mistaken example (x, y) contributes an addition $y\mathbf{x}$ (online) or $\frac{1}{n}y\mathbf{x}$ (batch) to the current weight vector
- For online perceptron, different orderings of the training examples can lead to different outputs
- Voted perceptron can handle non-linearly separable data, and is more robust to noise/outlier
- A very simple yet general algorithm that can be extended to handle much more complex scenarios